

Reviewing AI Output — Two Red Flags

*A take-home card for checking anything an AI drafts for you. “Verify everything” is hard to act on; these two tells show you **where to dig first**.*




Figure 1: Two AI red flags — verbosity (“it won’t stop explaining”) and certainty (“it sounds too sure”) — point to: dig here and verify. Plausible is not correct.

The two flags

1. **Verbosity** — it’s over-explaining. Reasoning models especially pad, restate, and pile on caveats. Length is often hiding thin or wrong substance. **When it won’t stop explaining, slow down and check what it’s actually claiming.**
2. **Certainty** — it sounds *too* sure. Hedge-free, confident statements with no “it depends” are a tell, not a comfort. The most confident sentence is often the one to verify. **Overconfidence is where it’s most likely confidently wrong.**

Rule of thumb: *the more it over-explains and the more certain it sounds, the harder you check.*
Plausible is not correct.

What to do when a flag trips

- Trace the specific claim back to a source you trust (the real tool’s behavior, the data, the docs).
- Ask it to state its uncertainty: “What in this are you least sure about?”
- Cross-check with a second prompt or a second model.

Remember

The AI **advises or drafts**; you **decide and verify** (the Command Rule). It’s wrong a meaningful fraction of the time, in *subtle* ways — these two flags are your fast triage for finding it.